

UNIVERSITY OF HOUSTON

INTRODUCTION TO COMPUTER NETWORKS

COSC 6377

Midterm Review

Author

K.M. HOURANI

Based on Notes By

Dr. Omprakash GNAWALI

March 3, 2021

Contents

1	Intro	5
1.1	The Internet	5
1.2	Packet vs. Circuit Switching	5
1.2.1	Circuit Switching	5
1.2.2	Some Circuit Switching Techniques	6
1.2.3	Packet Switching	6
1.2.4	Summary	6
1.3	Describing a Network	6
1.3.1	Throughput	7
1.3.2	Latency	7
1.3.2.1	Relation between Latency and Throughput	7
1.3.3	Reliability	7
1.4	Protocols	7
1.5	Network Protocols	8
1.5.1	Protocols and Standards	8
1.5.1.1	Protocol Layers	8
1.6	Encapsulation	8
2	Network Applications and Socket Programming	11
2.1	Network Applications	11
2.1.1	Inter-Application Communication	11
2.1.2	Application Protocols	11
2.1.3	Network Time Service	12
2.1.3.1	Protocol Timing Diagram	12
2.1.3.2	Cloud-based File Backup Application	12
2.2	Socket Programming	13
2.2.1	Using TCP/IP	13
2.2.2	System Calls	13
2.2.3	File Descriptors	13
2.2.4	Error Returns	13
2.2.5	Some operations on File Descriptors	14
2.2.6	Sockets: Communication Between Machines	14
2.2.7	System calls for using TCP	14
2.2.8	Socket Naming	14
2.2.9	Socket Address Structures	15
2.2.10	Dealing with Address Types	15
2.2.11	Client Skeleton (IPv4)	15
2.2.12	Server Skeleton (IPv4)	15
2.2.13	Looking up socket address with <code>getaddrinfo</code>	16
2.2.14	<code>getaddrinfo()</code> [RFC3493]	16
2.2.15	EOF in more detail	16
2.2.16	Using UDP	17
2.2.17	Serving Multiple Clients	17

2.2.18	Threads	17
2.2.19	Non-blocking I/O	17
2.2.20	How do you know when to read/write?	18
2.2.21	Event-driven servers	18


Chapter 1

Intro

1.1 The Internet

- Collection of nodes, wired and wireless technology connecting these nodes, applications and services
- Types of nodes
 - Desktops and Laptop
 - Servers
 - TV/Refrigerator
 - Cellphones
- Goal: Connect all the nodes to each other
- Solutions
 - $\binom{n}{2} = \mathcal{O}(n^2)$ cables
 - Sharing the links
 - * Circuit Switching
 - * Packet Switching
- Packet
 - Collection of bits to transfer across a network
 - Think: envelope and its contents
- Circuit
 - Pre-allocated path/resource


1.2 Packet vs. Circuit Switching



1.2.1 Circuit Switching

- Setup the connection or resource

- Schedule (e.g., TDMA)
- State in the network



Time	Circuit
$T, 3T, 5T, \dots$	$A - D$
$2T, 4T, 6T, \dots$	$B - C$

- Natural for predictable data rates
- Can guarantee certain level of services
- Can be inefficient for many applications

1.2.2 Some Circuit Switching Techniques

- Time
 - Reserve to use the link at a given schedule
 - Read: https://en.wikipedia.org/wiki/Time-division_multiplexing
- Frequency
 - Reserve to use certain frequencies (channel)
 - Read: https://en.wikipedia.org/wiki/Frequency-division_multiplexing

1.2.3 Packet Switching

- Wire is selected for each packet
- No network **state**
- Supports unpredictable/bursty traffic pattern
- Higher link utilization
- No guarantees but good enough for most applications

https://en.wikipedia.org/wiki/Packet_switching

1.2.4 Summary

- Packet Switching
 - Plus: more sharing (more efficient)
 - Minus: no service guarantee
- Circuit Switching
 - Plus: service guarantee
 - Minus: less sharing (less efficient)
- Every day examples
 - Road network

1.3 Describing a Network

- How to describe how well a network is working?
 - Metrics
- Performance metrics
 - Throughput
 - Latency
 - Reliability

1.3.1 Throughput

- How many bytes can we send through in a given time?
 - Bytes per second
 - How many bits/s in kbps?
 - Read: https://en.wikipedia.org/wiki/Data-rate_units
- Useful bytes transferred vs. overhead
 - Goodput
 - Everyday example: car vs. passenger

<https://en.wikipedia.org/wiki/Throughput>

1.3.2 Latency

- How long does it take for one bit to travel from one end to the other end?
 - ms, s, minutes, etc.
- Typical latencies
 - Speed of light
 - Why is web browsing latency in seconds?

1.3.2.1 Relation between Latency and Throughput

- Characterize the latency and throughput of
 - Oil Tanker –
 - Aircraft –
 - Car –
 - Tractor Trailer –
- Which metrics matter most for these applications?
 - Netflix
 - Skype
 - Amazon
 - Facebook

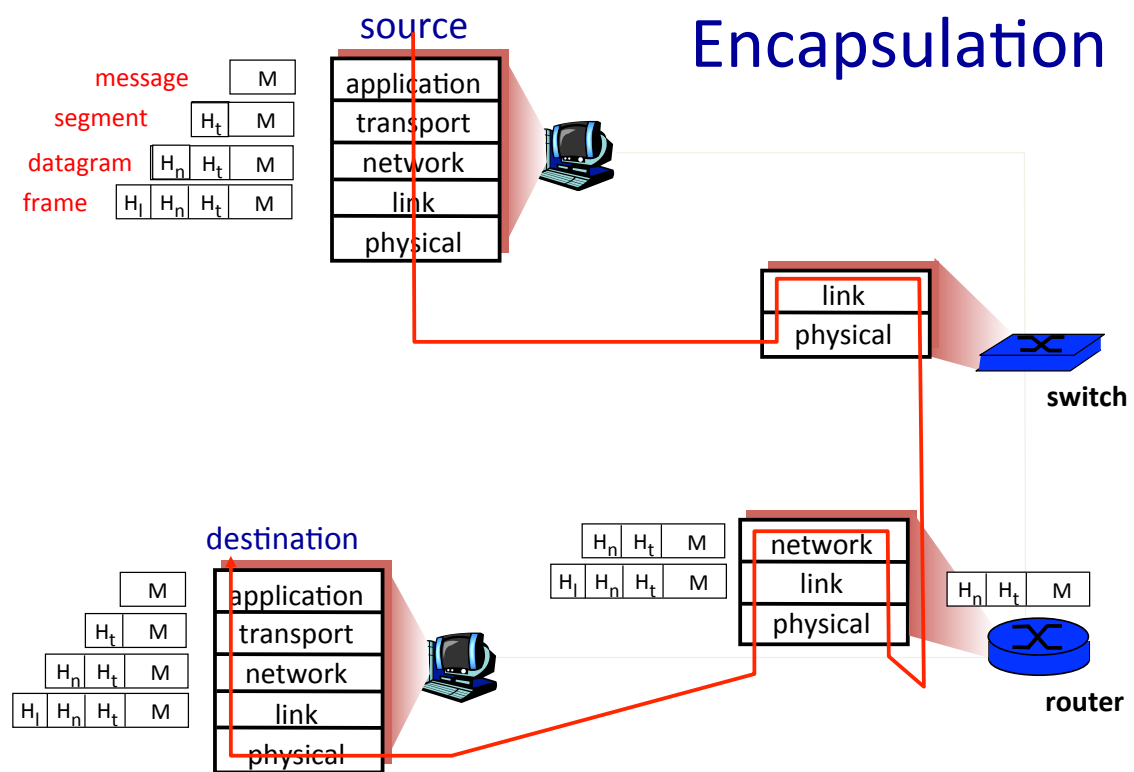
1.3.3 Reliability

- How often does a network fail?
- How often do packets drop?
 - Damage (corruption)
 - Drops in the queues
- How persistent are failures?
- Typical metrics
 - uptime percentage
 - packet or bit loss rate

1.4 Protocols

- Agreed-upon rules, format, and meaning for message exchange
- Let's examine this sequence:
 - Hello
 - How are you?
 - Fine.


https://en.wikipedia.org/wiki/Communication_protocol



Chapter 2

Network Applications and Socket Programming


2.1 Network Applications



<https://en.wikipedia.org/wiki/Peer-to-peer>

2.1.1 Inter-Application Communication


- Need a way to send and receive messages
- Inter-process communication
- Need naming, routing, transport
- Transport using TCP and UDP
 - On top of IP



2.1.2 Application Protocols


- Messages between processes, typically encapsulated within TCP or UDP
- Need agreement between
 - Sending process
 - Receiving process

2.1.3 Network Time Service




2.1.3.1 Protocol Timing Diagram

Protocol Timing Diagram



2.1.3.2 Cloud-based File Backup Application

- Client-server or peer-to-peer?
- Where do the applications run?
- Who/how to run these applications?
- What messages are exchanged?




2.2 Socket Programming

2.2.1 Using TCP/IP

- How can applications use the network?
- *Sockets* API
 - Originall from BS, widely implemented (*BSD, Linux, Mac OS X, Windows, ...)
 - Higher-level APIs build on them
- After basic setup, much like files

One could test network protocols with read/write on a file



2.2.2 System Calls

- Problem: how to access resources other than the CPU
 - Disk, netowrk, terminal, other processes
 - CPU prohibits instructions that would access devices
 - Only privileged OS kernel can access devices
- Kernel supplies well-defined system call interface
 - Applications request I/O oeprations through syscalls
 - Set up syscall arguments and trap to kernel
 - Kernel performs operation and returns result
- Higher-level functions built on syscall interface
 - `printf`, `scanf`, `gets`, all user-level code

2.2.3 File Descriptors

- Most I/O in Unix done through *file descriptors*
 - Integer *handles* to per-process table in kernel
- `int open(char *path, int flags, ...);`
- Returns file descriptor, used for all I/O to file

https://en.wikipedia.org/wiki/File_descriptor

2.2.4 Error Returns

- What if `open` fails? Return `-1` (invalid file descriptor)
- Most system calls return `-1` on failure
 - Specific type of error in gobal `int errno`

- `#include <sys/errno.h>` for possible values
 - `2` = `ENOENT` “no such file or directory”
 - `13` = `EACCES` “permission denied”

2.2.5 Some operations on File Descriptors

- `ssize_t read(int fd, void* buf, int nbytes);`
 - Returns number of bytes read
 - Returns `0` bytes at end of file, or `-1` on error
- `ssize_t write(int fd, void* buf, int nbytes);`
 - Returns number of bytes written, `-1` on error
- `off_t lseek(int fd, off_t offset, int whence);`
 - `whence`: `SEEK_SET`, `SEEK_CUR`, `SEEK_END`
 - returns new offset, or `-1` on error
- `int close(int fd);`

2.2.6 Sockets: Communication Between Machines

- Network sockets are file descriptors too
- Datagram sockets: unreliable message delivery
 - With IP, gives you UDP
 - Send atomic messages, which may be reordered or lost
 - Special system calls to read/write: `send/recv`
- Stream sockets: bi-directional pipes
 - With IP, gives you TCP
 - Bytes written on one end read on another
 - Reads may not return full amount requested, must reread

2.2.7 System calls for using TCP

<u>Client</u>	<u>Server</u>
1.	<code>socket</code> – make socket
2.	<code>bind</code> – assign address, port
3.	<code>listen</code> – listen for clients
4. <code>socket</code> – make socket	
5. <code>bind</code> – assign address ¹	
6. <code>connect</code> – connect to listening socket	
7.	<code>accept</code> – accept connection

2.2.8 Socket Naming

- Naming of TCP and UDP communication endpoints
 - IP address specifies host (129.7.240.18)
 - 16-bit port number demultiplexes within host
 - Well-known services listen on standard ports (e.g. `ssh` – 22, `http` – 8, see `/etc/services` for list)
 - Clients connect from arbitrary ports to well-known ports
- A connection is named by 5 components
 - Protocol, local IP, local port, remote IP, remote port
 - TCP requires connected sockets, but not UDP

¹This call to `bind` is optional, `connect` can choose address and port

2.2.9 Socket Address Structures

- Socket interface supports multiple network types
- Most calls take a generic `sockaddr`:

```
struct sockaddr {
    uint16_t sa_family; /* address family */
    char      sa_data[14]; /* protocol-specific addr */
};
```

- e.g. `int connect(int s, struct sockaddr* srv, socklen_t addrlen);`
- Cast `sockaddr*` from protocol-specific struct, e.g.

```
struct addr_in {
    short    sin_family; /* = AF_INET */
    u_short  sin_port; /* = htons (PORT) */
    struct   in_addr sin_addr; /* 32-bit IPV4 addr */
    char     in_zero[8];
};
```

2.2.10 Dealing with Address Types

- All values in network byte order (Big Endian)
 - `htonl()`, `htons()`: host to network, 32 and 16 bits
 - `ntohl()`, `ntohs()`: network to host, 32 and 16 bits
 - **Remember to always convert!**
- All address types begin with family
 - `sa_family` in `sockaddr` tells you the actual type
- Not all addresses are the same size
 - e.g. `struct sockaddr_in6` is typically 28 bytes, yet generic `struct sockaddr` is only 16 bytes
 - so most calls require passing around socket length
 - new `sockaddr_storage` is big enough

2.2.11 Client Skeleton (IPv4)

```
struct sockaddr_in {
    short    sin_family; /* = AF_INET */
    u_short  sin_port; /* = htons (PORT) */
    struct   in_addr sin_addr;
    char     sin_zero[8];
} sin;

int s = socket (AF_INET, SOCK_STREAM, 0);
memset(&sin, sizeof(sin), 0);
sin.sin_family = AF_INET;
sin.sin_port = htons(13); /* daytime port */
sin.sin_addr.s_addr = htonl(IP_ADDRESS);
connect(s, (sockaddr*)&sin, sizeof(sin));
while ((n = read(s, buf, sizeof(buf))) > 0) {
    write(1, buf, n);
}
```

2.2.12 Server Skeleton (IPv4)

```
int s = socket(AF_INET, SOCK_STREAM, 0);
struct sockaddr_in sin;
```

```

memset(&sin, sizeof(sin), 0);
sin.sin_family = AF_INET;
sin.sin_port = htons(9999);
sin.sin_addr.s_addr = htonl(INADDR_ANY);
bind(s, (struct sockaddr*)&sin, sizeof(sin));
listen(s, 5);
while (true) {
    socklen_t len = sizeof(sin);
    int cfd = accept(s, (struct sockaddr*)&sin, &len);
    /* cfd is new connection; you never read/write s */
    do_something_with(cfd);
    close(cfd);
}

```

2.2.13 Looking up socket address with getaddrinfo

```

struct addrinfo hints, *ai;
int err;
memset(&hints, 0, sizeof(hints));
hints.ai_family = AF_UNSPEC; /* or AF_INET or AF_INET6 */
hints.ai_socktype = SOCK_STREAM; /* or SOCK_DGRAM for UDP */

err = getaddrinfo("www.brown.edu", "http", &hints, &ai);
if (err) {
    fprintf(stderr, "%s\n", gai_strerror(err));
} else {
    /* ai->ai_family = address type (AF_INET or AF_INET6) */
    /* ai->ai_addr = actual address cast to (sockaddr *) */
    /* ai->ai_addrlen = length of actual address */
    freeaddrinfo(ai); /* must free when done! */
}

```

2.2.14 getaddrinfo()[RFC3493]

- Protocol-independent node name to address translation
 - Can specify port as a service name or number
 - May return multiple addresses
 - You must free the structure with `freeaddrinfo`
- Other useful functions to know about
 - `getnameinfo` – lookup hostname based on address
 - `inet_ntop` – convert IPv4 or 6 address to printable
 - `inet_prton` – convert string to IPv4 or 6 address

2.2.15 EOF in more detail

- What happens at the end of store?
 - Server receives EOF, renames file, responds OK
 - Client reads OK, *after* sending EOF: didn't close fd
- `int shutdown(int fd, int how);`
 - Shuts down a socket without closing the file descriptor
 - how: 0 = read, 1 = write, 2 = both
 - Note 1: applies to *socket*, not descriptor, so copies of descriptor (through fork or dup) affected
 - Note 2: with TCP, can't detect if other side shuts down for reading

2.2.16 Using UDP

- Call socket with `SOCK_DGRAM`, bind as before
- New calls for sending/receiving individual packets
 - `sendto(int s, const void* msg, int len, int flags, const struct sockaddr* to, socklen_t tolen);`
 - `recvfrom(int s, void* buf, int len, int flags, struct sockaddr *from, socklen_t* fromlen);`
 - Must send/get peer address with each packet
- Can use UDP in connected mode (why?)
 - connect assigns remote address
 - `send/recv` syscalls, like `sendto/recvfrom`, without last two arguments

2.2.17 Serving Multiple Clients

- A server may block when talking to a client
 - Read or write of a socket connected to a slow client can block
 - Server may be busy with CPU
 - Server might be blocked waiting for disk I/O
- Concurrency through multiple processes
 - Accept, fork, close in parent; child services request
- Advantages of one process per client
 - Doesn't block on slow clients
 - May use multiple cores
 - Can keep disk queues full for disk-heavy workloads

2.2.18 Threads

- One process per client has disadvantages:
 - High overhead – `fork + exit` $\approx 100\mu\text{sec}$
 - Hard to share state across clients
 - Maximum number of processes limited
- Can use threads for concurrency
 - Data races and deadlocks make programming tricky
 - Must allocate one stack per request
 - Many thread implementations block on some I/O or have heavy thread-switch overhead

Rough equivalents to `fork()`, `waitpid()`, `exit()`, `kill()`, plus locking primitives.

2.2.19 Non-blocking I/O

- `fcntl` sets `O_NONBLOCK` flag on descriptor

```
int n;
if ((n = fcntl(s, F_GETFL)) >= 0) {
    fcntl(s, F_SETFL, n | O_NONBLOCK);
}
```

- Non-blocking semantics of system calls:
 - read immediately returns `-1` with `errno` `EAGAIN` if no data
 - write may not write all data, or may return `EAGAIN`
 - connect may fail with `EINPROGRESS` (or may succeed, or may fail with a real error like `ECONNREFUSED`)
 - accept may fail with `EAGAIN` or `EWOULDBLOCK` if no connections present to be accepted

2.2.20 How do you know when to read/write?

```
struct timeval {
    long tv_sec; /* seconds */
    long tv_usec; /* and microseconds */
};
int select(int nfds, fd_set* readfds, fd_set* writefds,
           fd_set* exceptfds, struct timeval* timeout);
FD_SET(fd, &fdset);
FD_CLR(fd, &fdset);
FD_ISSET(fd, &fdset);
FD_ZERO(&fdset);
```

- Entire program runs in an *event loop*

2.2.21 Event-driven servers

- Quite different from processes/threads
 - Race conditions, deadlocks rare
 - Often more efficient
- But...
 - Unusual programming model
 - Sometimes difficult to avoid blocking
 - Scaling to more CPUs is more complex